

Operating Systems


Recitation 8

Plan

- Thread synchronization
 - Lock & event in Linux
 - Mutex & Condition variable
 - What happens “under the hood”

Reminder

- Threads get scheduled by scheduler in kernel
- Preemptive multitasking:
 - OS decides when a thread will get its CPU time slot
- Context-switch without warning



Reminder

```
static int counter = 1;
```


```
int next_counter(void) {  
    return counter++;  
}
```

**OK for single thread,
not for concurrent threads**

What It Means

```
static int counter = 1;
```

```
int next_counter(void) {  
    return counter++;  
}
```



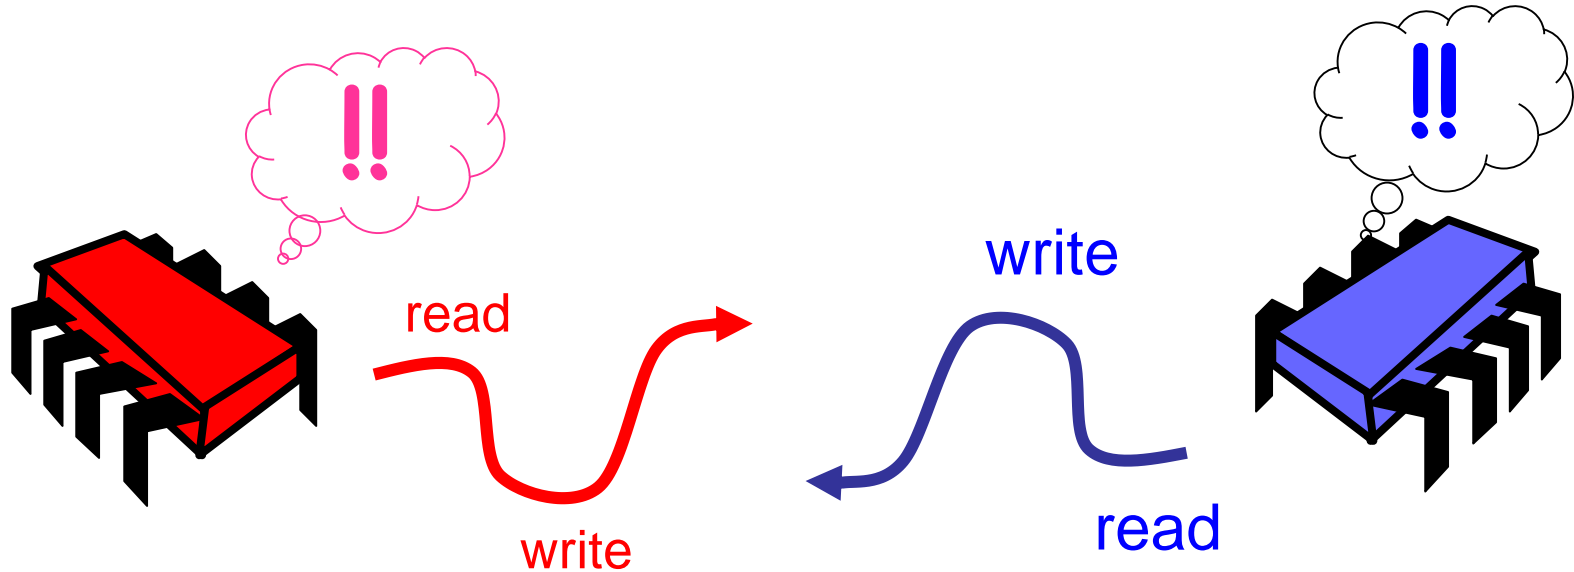
```
temp = value;  
value = temp + 1;  
return temp;
```

Race condition

thread 1	thread 2	value
temp = value		0
	temp = value	0
	temp = temp + 1	0
	Value = temp	1
temp = temp + 1		1
value = temp		1 (not 2!)

- Result depends on the timing execution of the code
- Can get different result every time!

Is this problem inherent?



What we have here is a **critical section**

If we could only “glue” reads and writes together...i.e. make them atomic...

Hardware Solution

- Read-Modify-Write instructions
- Atomic operations at **CPU level!**
 - Supported in instruction set
 - Example: intel x86 supports *atom_inc*, *atom_dec*, *atom_add*, *atom_sub*
- Built-in functions for atomic memory access
- type `__sync_fetch_and_add` (type *ptr, type value, ...)
- *Code example 1*

Linux Synchronization

- Main mechanism implemented in Linux for user-level synchronization
 - **Mutex** (lock)
 - **Condition variable** (signal)

Mutex

- Mutual exclusion
- Only **one** thread can hold it **lock()**ed
 - Others trying to **lock()** block until owner decides to free

- *Example:*

lock

```
/* do some critical section code */
```

unlock

Mutex Formal Requirements

- Mutually exclusive
 - Only **one** thread can hold it lock()ed – can't have two in same critical section
- Deadlock-free
 - **Some** thread eventually enters critical section
- Starvation-free
 - **Each** thread eventually enters critical section

It's all in your head!

- Always remember when programming with mutexes:
it's a logical concept
- The protection of variables and code sections exists only *in your head*
- If you don't **consistently** protect shared variables/critical code with a mutex, bad things will happen
- OS only provides the mechanism - you are the user!



Mutex API

- Creation:

```
int pthread_mutex_init(  
    pthread_mutex_t *mutex,  
    const pthread_mutex_attr_t *mutexattr);
```

- Destruction:

```
int pthread_mutex_destroy(  
    pthread_mutex_t *mutex);
```

Mutex API - Example

```
pthread_mutex_t lock;
```

```
int main() {  
    if (pthread_mutex_init(&lock, NULL) != 0) {  
        perror("mutex init failed\n");  
        return 1;  
    }  
    /* ... code here ... */  
    pthread_mutex_destroy(&lock);  
}
```

Locking

Lock:

`int pthread_mutex_lock(pthread_mutex_t *mutex);`

- Acquire lock, block until acquired

`int pthread_mutex_trylock(pthread_mutex_t *mutex);`

- Acquire lock once, fail if already locked

Unlock:

`int pthread_mutex_unlock(pthread_mutex_t *mutex);`

- Release locked mutex

- ***Code example 2***

Under the Hood

- Known algorithms: **Peterson's, bakery**
 - Use a lot of memory ($O(t)$ for t threads)
- Another suggestion:
 - Shared global variable acts as a 'lock'
 - Initially 'unlocked'
 - `int mutex = 0;`
 - Before entering critical section, a task 'locks' the mutex
 - `mutex = 1;`
 - When done with critical section, 'unlocks' the mutex
 - `mutex = 0;`
 - While "locked", no other task can enter critical section

Under the Hood

```
void init() {  
    flag = 0;  
}  
void lock() {  
    // busy-wait, loop until value is 0 → unlocked!  
    while (flag == 1);  
    flag = 1;  
}  
void unlock() {  
    flag = 0;  
}
```

What's the problem?

Under the hood

- Special mutex variable needs to be accessed atomically
- Reasonable solution - hardware support
- One example (from the past):

testandset <address>, rnew, rold

- Special atomic operation

```
int TestAndSet(int *lock, int new) {  
    int old = *lock; // save old value of &lock in memory  
    *lock = new;      // set new value  
    return old;       // return old value  
}
```

Test-and-Set Implementation

```
void init() {  
    // 0 means lock is available, 1 means held by a thread  
    flag = 0;  
}  
  
void lock() {  
    // busy-wait (do nothing)  
    // exits loop only when old value is 0 == not locked!  
    while (TestAndSet(&flag, 1) == 1) ;  
}  
  
void unlock() {  
    flag = 0;  
}
```

Simple implementation

```
void init() {  
    // 0 means lock is available, 1 means held by a thread  
    flag = 0;  
}  
void lock() {  
    // busy-wait (do nothing)  
    // exits loop only when old value is 0 == not locked!  
    while (TestAndSet(&flag, 1) == 1) ;  
}  
void unlock() {  
    flag = 0;  
}
```

That's a LOT of *spinning*!
Too many time-slices
wasted by scheduler on
threads in hopeless loop

Also possibly *starvation*!
Doesn't ensure all threads
will eventually acquire lock!

Less naive implementation

- Add `yield()` instruction

```
void init() {  
    flag = 0;  
}  
void lock() {  
    while (TestAndSet(&flag, 1) == 1)  
        yield(); // give up CPU on lock failure  
}  
void unlock() {  
    flag = 0;  
}
```

Events

- Allow thread1 to inform thread2 on some event
 - Thread2 can sleep meanwhile
- Allow sync. access to sensitive shared resource
- Extension to mutex

Example: simple queue

- thread1 enqueues, thread2 dequeues
- Without sync. access:
 - Both threads may change data together
 - Thread1 insertion not safe (memory addresses...)
 - Thread2 won't know when to deq (memory addresses, polling...)
- Without events:
 - Possible deadlock?

Condition Variables (1)

- Allow thread to sleep-wait() on event

```
int pthread_cond_init(  
    pthread_cond_t *cond,  
    pthread_condattr_t *cond_attr);
```

```
int pthread_cond_destroy(pthread_cond_t *cond);
```

- Initialize/destroy condition variable object
 - cond_attr = NULL is default
- Destroy fails if threads are waiting

Condition Variables (2)

```
int pthread_cond_wait(  
    pthread_cond_t *cond,  
    pthread_mutex_t *mutex);
```

- Wait() on condition variable
- **Must have mutex already locked!**
- On success releases mutex and puts thread to sleep
- Several threads can wait()
 - But only one wakes up...

Condition Variables (3)

```
int pthread_cond_signal(pthread_cond_t *cond);
```

- Signal a single wait()ing thread to wake up
- Choice of awakened thread is arbitrary
- Notice – **no mutex**

Back to queue example

```
item dequeue() {  
    pthread_mutex_lock(&qlock);  
    while <queue is empty>  
        pthread_cond_wait(&notEmpty,&qlock);  
    /* ... remove item from queue ... */  
    pthread_mutex_unlock(&qlock);  
    /* .. return removed item */  
}
```

Why **while**?

Back to queue example

```
pthread_mutex_t qlock;  
pthread_cond_t notEmpty;  
/* ... initialization code ... */  
void enqueue(item x) {  
    pthread_mutex_lock(&qlock);  
    /* ... add x to queue ... */  
    pthread_cond_signal(&notEmpty);  
    pthread_mutex_unlock(&qlock);  
}
```